



Chapter 7

AVL Trees, B-Trees

Data Structures and Algorithms

LE Thanh Sach

*Faculty of Computer Science and Engineering
University of Technology, VNU-HCM*

- **L.O.3.1** - Depict the following concepts: binary tree, complete binary tree, balanced binary tree, AVL tree, multi-way tree, etc.
- **L.O.3.2** - Describe the storage structure for tree structures using pseudocode.
- **L.O.3.3** - List necessary methods supplied for tree structures, and describe them using pseudocode.
- **L.O.3.4** - Identify the importance of “balanced” feature in tree structures and give examples to demonstrate it.
- **L.O.3.5** - Identify cases in which AVL tree and B-tree are unbalanced, and demonstrate methods to resolve all the cases step-by-step using figures.



Outcomes

- **L.O.3.6** - Implement binary tree and AVL tree using C/C++.
- **L.O.3.7** - Use binary tree and AVL tree to solve problems in real-life, especially related to searching techniques.
- **L.O.3.8** - Analyze the complexity and develop experiment (program) to evaluate methods supplied for tree structures.
- **L.O.8.4** - Develop recursive implementations for methods supplied for the following structures: list, tree, heap, searching, and graphs.
- **L.O.1.2** - Analyze algorithms and use Big-O notation to characterize the computational complexity of algorithms composed by using the following control structures: sequence, branching, and iteration (not recursion).



Contents

- ➊ AVL Tree Concepts
- ➋ AVL Balance
- ➌ AVL Tree Operations
- ➍ Multiway Trees
- ➎ B-Trees





AVL Tree Concepts



Definition

AVL Tree is:

- A Binary Search Tree,
- in which the heights of the left and right subtrees of the root differ by at most 1, and
- the left and right subtrees are again AVL trees.

Discovered by G.M. [Adel'son-Vel'skii](#) and E.M. [Landis](#) in 1962.

[AVL Tree](#) is a Binary Search Tree that is balanced tree.

AVL Tree Concepts

AVL Balance

AVL Tree
Operations

Multiway Trees

B-Trees



A binary tree is an **AVL Tree** if

- **Each node satisfies BST property:** key of the node is greater than the key of each node in its left subtree and is smaller than or equals to the key of each node in its right subtree.
- **Each node satisfies balanced tree property:** the difference between the heights of the left subtree and right subtree of the node does not exceed one.

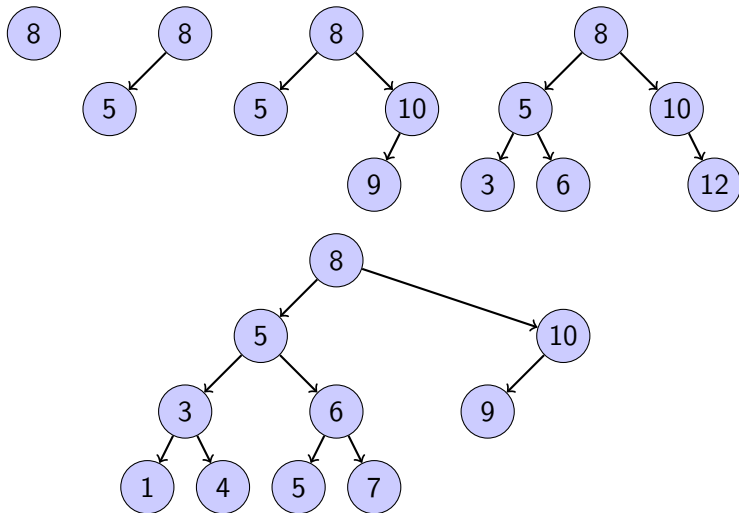


Balance factor

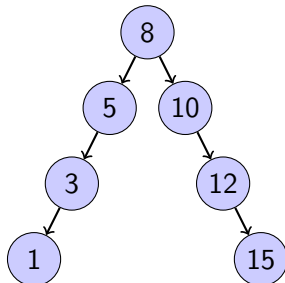



- left_higher (LH): $H_L = H_R + 1$
- equal_height (EH): $H_L = H_R$
- right_higher (RH): $H_R = H_L + 1$

(H_L, H_R : the heights of left and right subtrees)

AVL Trees



Non-AVL Trees



Why AVL Trees?

- When data elements are inserted in a BST in sorted order: 1, 2, 3, ...
BST becomes a degenerate tree.
Search operation takes $O(n)$, which is inefficient.
- It is possible that after a number of insert and delete operations, a binary tree may become unbalanced and increase in height.
- AVL trees ensure that the complexity of search is $O(\log_2 n)$.





AVL Balance

Balancing Trees

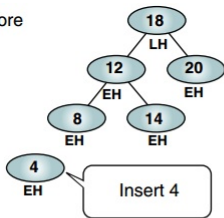
- When we insert a node into a tree or delete a node from a tree, the resulting tree may be unbalanced.
→ **rebalance the tree.**
- Four unbalanced tree cases:
 - **left of left**: a subtree of a tree that is left high has also become left high;
 - **right of right**: a subtree of a tree that is right high has also become right high;
 - **right of left**: a subtree of a tree that is left high has become right high;
 - **left of right**: a subtree of a tree that is right high has become left high;




Unbalanced tree cases



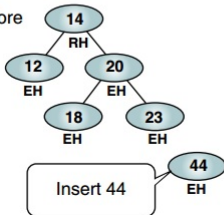
Before

**(a) Case 1: left of left**Out of balance
at 18

After




Before

**(b) Case 2: right of right**

(Source: Data Structures - A Pseudocode Approach with C++)

Out of balance
at 14


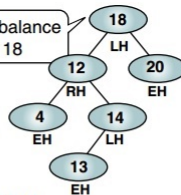
After



Unbalanced tree cases



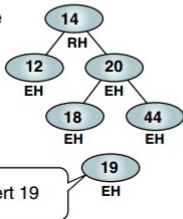
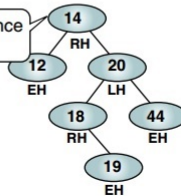
Before

Out of balance
at 18

After

(c) Case 3: right of left

Before

Out of balance
at 14

After

(d) Case 4: left of right

(Source: Data Structures - A Pseudocode Approach with C++)

Rotate Right

Algorithm rotateRight(ref root <pointer>)
Exchanges pointers to rotate the tree right.

Pre: root is pointer to tree to be rotated

Post: node rotated and root updated

tempPtr = root->left

root->left = tempPtr->right

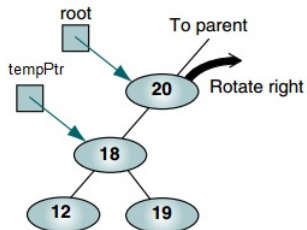
tempPtr->right = root

Return tempPtr

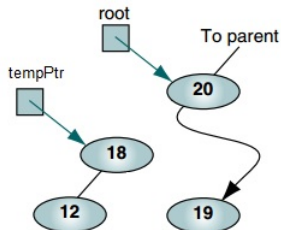
End rotateRight



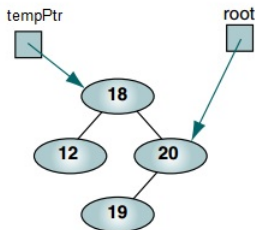
Rotate Right




(a) After Step 1



(b) After Step 2



(c) After Step 3



(d) At end

(Source: Data Structures - A Pseudocode Approach with C++)



Rotate Left

Algorithm rotateLeft(ref root <pointer>)
Exchanges pointers to rotate the tree left.

Pre: root is pointer to tree to be rotated

Post: node rotated and root updated

```
tempPtr = root->right  
root->right = tempPtr->left  
tempPtr->left = root  
Return tempPtr
```

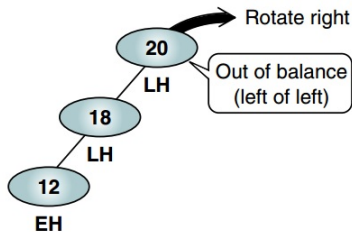
End rotateLeft



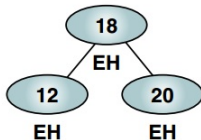
Balancing Trees - Case 1: Left of Left

Out of balance condition created by a left high subtree of a left high tree

→ balance the tree by rotating the out of balance node to the right.



(a1) After inserting 12




(a2) After rotation

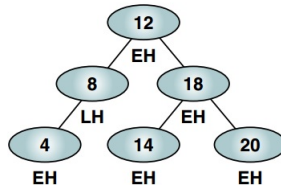
(Source: Data Structures - A Pseudocode Approach with C++)



Balancing Trees - Case 1: Left of Left



(b1) After inserting 4



(b2) After rotation


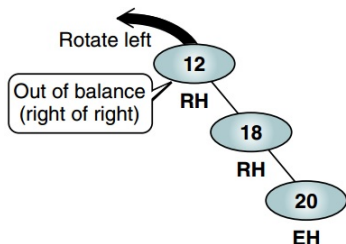
(Source: Data Structures - A Pseudocode Approach with C++)



Balancing Trees - Case 2: Right of Right

Out of balance condition created by a right high subtree of a right high tree

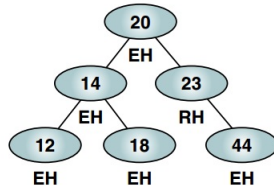

→ balance the tree by rotating the out of balance node to the left.



(Source: Data Structures - A Pseudocode Approach with C++)



Balancing Trees - Case 2: Right of Right



(Source: Data Structures - A Pseudocode Approach with C++)




Balancing Trees - Case 3: Right of Left

Out of balance condition created by a right high subtree of a left high tree

→ balance the tree by two steps:


- ① rotating the left subtree to the left;
- ② rotating the root to the right.




(Source: Data Structures - A Pseudocode Approach with C++)




Balancing Trees - Case 3: Right of Left



(b1) Original tree



(b2) After left rotation



(b3) After right rotation


(Source: Data Structures - A Pseudocode Approach with C++)

Balancing Trees - Case 4: Left of Right

Out of balance condition created by a left high subtree of a right high tree

→ balance the tree by two steps:

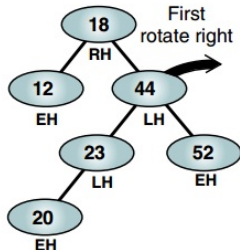
- ① rotating the right subtree to the right;
- ② rotating the root to the left.




(Source: Data Structures - A Pseudocode Approach with C++)



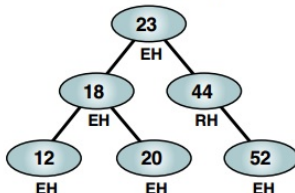
Balancing Trees - Case 4: Left of Right



(b1) Original tree



(b2) After right rotation



(b3) After left rotation

(Source: Data Structures - A Pseudocode Approach with C++)





AVL Tree Operations

AVL Tree Structure

```
node                                avlTree
  data <dataType>                  root <pointer>
  left <pointer>                   end avlTree
  right <pointer>
  balance <balance_factor>
end node
```

```
// General dataType:
dataType
  key <keyType>
  field1 <...>
  field2 <...>
  ...
  fieldn <...>
end dataType
```

Note: Array is not suitable for AVL Tree.



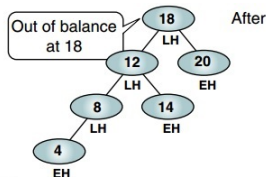

AVL Tree Operations

- Search and retrieval are the same for any binary tree.
- AVL Insert
- AVL Delete

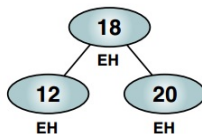





AVL Insert

- Insert can make an out of balance condition.



- Otherwise, some inserts can make an automatic balancing.



AVL Insert Algorithm

Algorithm AVLInsert(ref root <pointer>, val
newPtr <pointer>, ref taller <boolean>)
Using recursion, insert a node into an AVL tree.

Pre: root is a pointer to first node in AVL
tree/subtree

newPtr is a pointer to new node to be inserted

Post: taller is a Boolean: true indicating
the subtree height has increased, false
indicating same height

Return root returned recursively up the tree



AVL Insert Algorithm

```
// Insert at root  
if root null then  
    |   root = newPtr  
    |   taller = true  
    |   return root  
end
```



AVL Insert Algorithm

```
if newPtr->data.key < root->data.key then  
    root->left = AVLInsert(root->left, newPtr,  
        taller)  
    // Left subtree is taller  
    if taller then  
        if root is LH then  
            | root = leftBalance(root, taller)  
        else if root is EH then  
            | root->balance = LH  
        else  
            | root->balance = EH  
            | taller = false  
        end  
    end  
end
```



AVL Insert Algorithm

else

 root->right = AVLInsert(root->right, newPtr,
 taller)

// Right subtree is taller

if taller then

if root is LH then

 root->balance = EH

 taller = false

else if root is EH then

 root->balance = RH

else

 root = rightBalance(root, taller)

end

end

end

return root

End AVLInsert



AVL Left Balance Algorithm

Algorithm leftBalance(ref root <pointer>, ref taller <boolean>)

This algorithm is entered when the left subtree is higher than the right subtree.

Pre: root is a pointer to the root of the [sub]tree
taller is true

Post: root has been updated (if necessary)
taller has been updated



AVL Left Balance Algorithm

leftTree = root->left

// Case 1: Left of left. Single rotation right.

if leftTree is LH **then**

 root = rotateRight(root)

 root->balance = EH

 leftTree->balance = EH

 taller = false



AVL Left Balance Algorithm

// Case 2: Right of Left. Double rotation required.

else

rightTree = leftTree->right

if rightTree->balance = LH **then**

 root->balance = RH

 leftTree->balance = EH

else if rightTree->balance = EH **then**

 leftTree->balance = EH

else

 root->balance = EH

 leftTree->balance = LH

end

rightTree->balance = EH

root->left = rotateLeft(leftTree)

root = rotateRight(root)

taller = false

end

return root

End leftBalance



AVL Right Balance Algorithm

Algorithm rightBalance(ref root <pointer>,
ref taller <boolean>)

This algorithm is entered when the right subtree is higher than the left subtree.

Pre: root is a pointer to the root of the
[sub]tree
taller is true

Post: root has been updated (if necessary)
taller has been updated



AVL Right Balance Algorithm

rightTree = root->right

// Case 1: Right of right. Single rotation left.

if *rightTree is RH* **then**

 root = rotateLeft(root)

 root->balance = EH

 rightTree->balance = EH

 taller = false



AVL Right Balance Algorithm

// Case 2: Left of Right. Double rotation required.

else

leftTree = rightTree->left

if leftTree->balance = RH **then**

 root->balance = LH

 rightTree->balance = EH

else if leftTree->balance = EH **then**

 rightTree->balance = EH

else

 root->balance = EH

 rightTree->balance = RH

end

leftTree->balance = EH

root->right = rotateRight(rightTree)

root = rotateLeft(root)

taller = false

end

return root

End rightBalance



AVL Delete Algorithm

The AVL delete follows the basic logic of the binary search tree delete with the addition of the logic to balance the tree. As with the insert logic, the balancing occurs as we back out of the tree.

Algorithm AVLDelete(ref root <pointer>, val deleteKey <key>, ref shorter <boolean>, ref success <boolean>)
This algorithm deletes a node from an AVL tree and rebalances if necessary.

Pre: root is a pointer to the root of the [sub]tree
deleteKey is the key of node to be deleted

Post: node deleted if found, tree unchanged if not found
shorter is true if subtree is shorter
success is true if deleted, false if not found

Return pointer to root of (potential) new subtree



AVL Delete Algorithm

```
if tree null then
    shorter = false
    success = false
    return null
end
if deleteKey < root->data.key then
    root->left = AVLDelete(root->left, deleteKey,
    shorter, success)
    if shorter then
        | root = deleteRightBalance(root, shorter)
    end
else if deleteKey > root->data.key then
    root->right = AVLDelete(root->right, deleteKey,
    shorter, success)
    if shorter then
        | root = deleteLeftBalance(root, shorter)
    end
```



AVL Delete Algorithm

// Delete node found – test for leaf node

else

deleteNode = root

if *no right subtree* **then**

newRoot = root->left

success = true

shorter = true

recycle(deleteNode)

return newRoot

else if *no left subtree* **then**

newRoot = root->right

success = true

shorter = true

recycle(deleteNode)

return newRoot



AVL Delete Algorithm

```
else
    // ... // Delete node has two subtrees
    else
        exchPtr = root->left
        while exchPtr->right not null do
            | exchPtr = exchPtr->right
        end
        root->data = exchPtr->data
        root->left = AVLDelete(root->left,
            exchPtr->data.key, shorter, success)
        if shorter then
            | root = deleteRightBalance(root, shorter)
        end
    end
end
end
Return root
End AVLDelete
```



Delete Right Balance

Algorithm deleteRightBalance(ref root <pointer>, ref shorter <boolean>)

The (sub)tree is shorter after a deletion on the left branch. Adjust the balance factors and if necessary balance the tree by rotating left.

Pre: tree is shorter

Post: balance factors updated and balance restored
root updated
shorter updated

if *root LH* **then**

 | root->balance = EH

else if *root EH* **then**

 | root->balance = RH

 | shorter = false



Delete Right Balance

else

rightTree = root->right

if rightTree LH **then**

leftTree = rightTree->left

if leftTree LH **then**

rightTree->balance = RH

root->balance = EH

else if leftTree EH **then**

root->balance = LH

rightTree->balance = EH

else

root->balance = LH

rightTree->balance = EH

end

leftTree->balance = EH

root->right = rotateRight(rightTree)

root = rotateLeft(root)



Delete Right Balance

```
else
    // ...
    else
        if rightTree not EH then
            root->balance = EH
            rightTree->balance = EH
        else
            root->balance = RH
            rightTree->balance = LH
            shorter = false
        end
        root = rotateLeft(root)
    end
end
return root
End deleteRightBalance
```



Delete Left Balance

Algorithm deleteLeftBalance(ref root <pointer>, ref shorter <boolean>)

The (sub)tree is shorter after a deletion on the right branch. Adjust the balance factors and if necessary balance the tree by rotating right.

Pre: tree is shorter

Post: balance factors updated and balance restored
root updated
shorter updated

if *root RH* **then**

 | root->balance = EH

else if *root EH* **then**

 | root->balance = LH

 | shorter = false



Delete Left Balance

else

leftTree = root->left

if *leftTree RH* **then**

rightTree = leftTree->right

if *rightTree RH* **then**

leftTree->balance = LH

root->balance = EH

else if *rightTree EH* **then**

root->balance = RH

leftTree->balance = EH

else

root->balance = RH

leftTree->balance = EH

end

rightTree->balance = EH

root->left = rotateLeft(leftTree)

root = rotateRight(root)



Delete Left Balance

```
else
    // ...
    else
        if leftTree not EH then
            root->balance = EH
            leftTree->balance = EH
        else
            root->balance = LH
            leftTree->balance = RH
            shorter = false
        end
        root = rotateRight(root)
    end
end
return root
End deleteLeftBalance
```





Multiway Trees

Multiway Trees

Tree whose outdegree is **not restricted to 2** while retaining the general properties of **binary search trees**.

AVL Trees, B-Trees

LE Thanh Sach



AVL Tree Concepts

AVL Balance

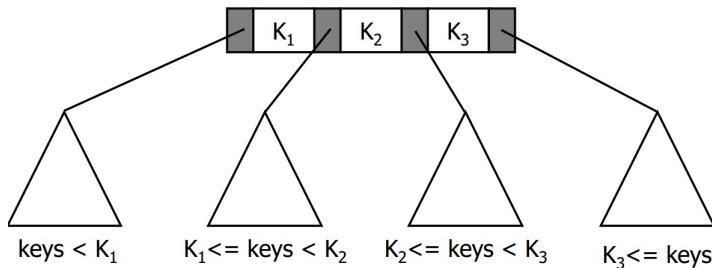
AVL Tree
Operations

Multiway Trees

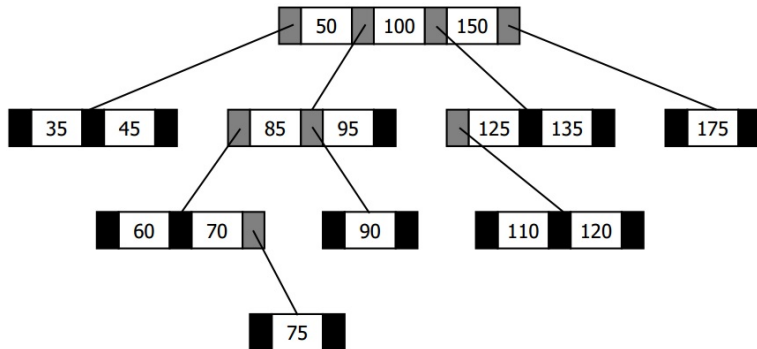
B-Trees

M-Way Search Trees


- Each node has $m - 1$ data entries and m subtree pointers.
- The key values in a subtree such that:
 - \geq the key of the left data entry
 - $<$ the key of the right data entry.




M-Way Search Trees



M-Way Node Structure



```
entry
  key <key type>
  data <data type>
  rightPtr <pointer>
end entry
```



```
node
  firstPtr <pointer>
  numEntries <integer>
  entries <array[1 .. m-1] of entry>
end node
```





B-Trees

- M-way trees are unbalanced.
- Bayer, R. & McCreight, E. (1970) created B-Trees.

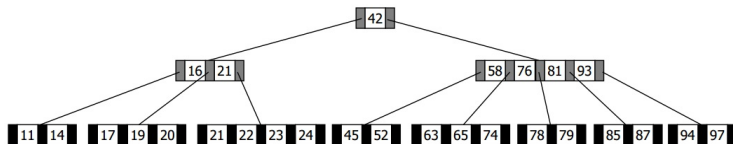


A B-tree is an m -way tree with the following additional properties ($m \geq 3$):

- The root is either a leaf or has at least 2 subtrees.
- All other nodes have at least $\lceil m/2 \rceil - 1$ entries.
- All leaf nodes are at the same level.



B-Trees



Hình: $m = 5$

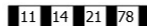


- Insert the new entry into a leaf node.
- If the leaf node is overflow, then split it and insert its median entry into its parent.

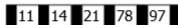


B-Tree Insertion

Insert 78, 21, 14, 11



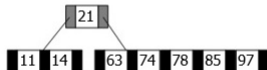
Insert 97



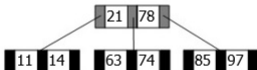
overflow



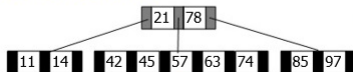
Insert 85, 74, 63



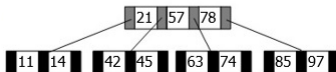
overflow



Insert 45, 42, 57





overflow





B-Tree Insertion

Insert 20, 16, 19



Insert 52, 30, 21



B-Tree Insertion

Algorithm BTreeInsert(ref root <pointer>, val data <record>)

Inserts data into B-tree. Equal keys placed on right branch.

Pre: root is a pointer to the B-tree. May be null.

Post: data inserted

Return pointer to B-tree root.

taller = insertNode(root, data, upEntry)

if taller **then**

// Tree has grown. Create new root.

 allocate(newPtr)

 newPtr->entries[1] = upEntry

 newPtr->firstPtr = root

 newPtr->numEntries = 1

 root = newPtr

end

return root

End BTreeInsert



B-Tree Insertion

Algorithm insertNode (ref root <pointer>, val data <record>, ref upEntry <entry>)

Recursively searches tree to locate leaf for data. If node overflow, inserts median key's data into parent.

Pre: root is a pointer to tree or subtree. May be null.

Post: data inserted

upEntry is overflow entry to be inserted into parent.

Return tree taller <boolean>.

if *root null* **then**

 upEntry.data = data

 upEntry.rightPtr = null

 taller = true



B-Tree Insertion

```
else
|   entryNdx = searchNode(root, data.key)
|   if entryNdx > 0 then
|       |   subTree = root->entries[entryNdx].rightPtr
|   else
|       |   subTree = root->firstPtr
|   end
|   taller = insertNode(subTree, data, upEntry)
|   if taller then
|       |   if node full then
|       |       |   splitNode(root, entryNdx, upEntry)
|       |       |   taller = true
|       |   else
|       |       |   insertEntry(root, entryNdx, upEntry)
|       |       |   taller = false
|       |       |   root->numEntries = root->numEntries + 1
|       |   end
|   end
|   end
end
return taller
End insertNode
```



B-Tree Insertion

Algorithm `searchNode(val nodePtr <pointer>, val target <key>)`

Search B-tree node for data entry containing key \leq target.

Pre: `nodePtr` is pointer to non-null node.

`target` is key to be located.

Return index to entry with key \leq target.

0 if key $<$ first entry in node

if `target < nodePtr->entry[1].data.key` **then**

`walker = 0`

else

`walker = nodePtr->numEntries`

while `target < nodePtr->entries[walker].data.key` **do**

`walker = walker - 1`

end

end

return `walker`

End `searchNode`



B-Tree Insertion

Algorithm splitNode(val node <pointer>, val entryNdx <index>, ref upEntry <entry>)

Node has overflowed. Split node. **No duplicate keys allowed.**

Pre: node is pointer to node that overflowed.

entryNdx contains index location of parent.

upEntry contains entry being inserted into split node.

Post: upEntry now contains entry to be inserted into parent.

minEntries = minimum number of entries

allocate (rightPtr)

// Build right subtree node

if entryNdx \leq minEntries **then**

 fromNdx = minEntries + 1

else



B-Tree Insertion

```
else
    | fromNdx = minEntries + 2
end
toNdx = 1
rightPtr->numEntries = node->numEntries - fromNdx
+ 1
while fromNdx <= node->numEntries do
    | rightPtr->entries[toNdx] = node->entries[fromNdx]
    | fromNdx = fromNdx + 1
    | toNdx = toNdx + 1
end
node->numEntries =
node->numEntries - rightPtr->numEntries
if entryNdx <= minEntries then
    | insertEntry(node, entryNdx, upEntry)
else
```



B-Tree Insertion

```
else
    insertEntry(rightPtr, entryNdx-minEntries,
upEntry)
    node->numEntries = node->numEntries- 1
    rightPtr->numEntries = rightPtr->numEntries + 1
end
// Build entry for parent
medianNdx = minEntries + 1
upEntry.data = node->entries[medianNdx].data
upEntry.rightPtr = rightPtr
rightPtr->firstPtr = node->entries[medianNdx].rightPtr
return
End splitNode
```



B-Tree Insertion

Algorithm insertEntry(val node <pointer>, val entryNdx <index>, val newEntry <entry>)

Inserts one entry into a node by shifting nodes to make room.

Pre: node is pointer to node to contain data.
entryNdx is index to location for new data.
newEntry contains data to be inserted.

Post: data has been inserted in sequence.

shifter = node->numEntries + 1

while shifter > entryNdx + 1 **do**

 node->entries[shifter] = node->entries[shifter - 1]
 shifter = shifter - 1

end

node->entries[shifter] = newEntry

node->numEntries = node->numEntries + 1

return

End insertEntry





B-Tree Deletion

- It must take place at a leaf node.
- If the data to be deleted are not in a leaf node, then replace that entry by the largest entry on its left subtree.


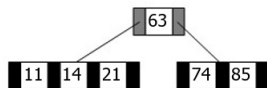


B-Tree Deletion

Delete 78





Delete 63





B-Tree Deletion

Delete 85



Delete 21



underflow
(node has fewer than the
min num of entries)



For each node to have sufficient number of entries:


- **Balance**: shift data among nodes.
- **Combine**: join data from nodes.




Balance

Borrow from right

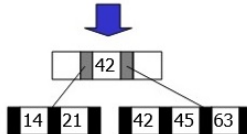
Original node



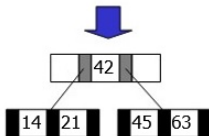
Rotate parent data down



Rotate data to parent



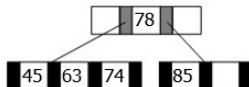
Shift entries left



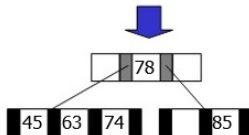
Balance

Borrow from left

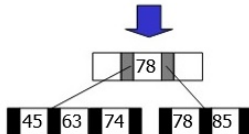
Original node



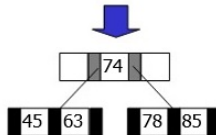
Shift entries right




Rotate parent data down



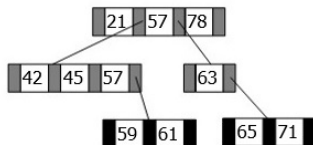
Rotate data up




Combine



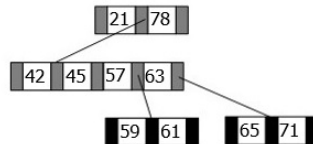
1. After underflow



2. After moving root to subtree



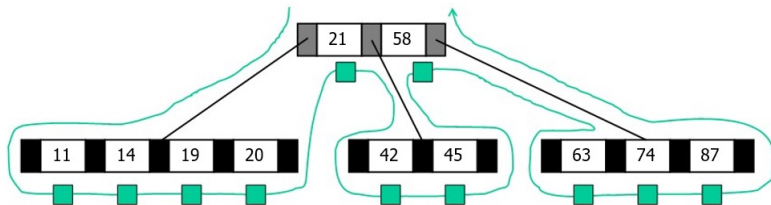
3. After moving right entries



4. After shifting root



B-Tree Traversal



B-Tree Traversal

Algorithm BTreeTraversal (val root <pointer>)

Processes tree using inorder traversal.

Pre: root is pointer to B-Tree.

Post: Every entry has been processed in order.

scanCount = 0

ptr = root->firstPtr

while scanCount <= root->numEntries **do**

if ptr not null **then**

 | BTreeTraversal(ptr)

end

 scanCount = scanCount + 1

if scanCount <= root->numEntries **then**

 | process (root->entries[scanCount].data)

 | ptr = root->entries[scanCount].rightPtr

end

end

return

End BTreeTraversal



B-Tree Search

Algorithm BTreeSearch(val root <pointer>, val target <key>, ref node <pointer>, ref entryNo <index>)
Recursively searches a B-tree for the target key.

Pre: root is pointer to a tree or subtree
target is the data to be located

Post:

if found – –

node is pointer to located node

entryNo is entry within node

if not found – –

node is null and entryNo is zero

Return found <boolean>



B-Tree Search

```
if target < first entry then
    |   return BTreeSearch (root->firstPtr, target, node,
    |   entryNo)
else
    |   entryNo = root->numEntries
    |   while target < root->entries[entryNo].data.key do
    |       |   entryNo = entryNo - 1
    |   end
    |   if target = root->entries[entryNo].data.key then
    |       |   found = true
    |       |   node = root
    |   else
    |       |   return BTreeSearch
    |       |   (root->entries[entryNo].rightPtr, target, node,
    |       |   entryNo)
    |   end
end
return found
End BTreeSearch
```



- **B*Tree**: the minimum number of (used) entries is two thirds.
- **B+Tree**:
 - Each data entry must be represented at the leaf level.
 - Each leaf node has one additional pointer to move to the next leaf node.

